

Influence Finder

An internal targeting recommendation tool for BPI campaign strategists

Michael Gioia · May 2026

1. The Opportunity

BPI strategists are walking into client meetings with a hard problem: given a specific campaign goal and target audience, which journalists, outlets, and creators are actually worth pursuing? Right now, that answer lives in experience and instinct — or in an ad hoc export from a data platform that nobody has time to properly parse.

The data to answer this question already exists. BPI has access to a graph dataset of 11,500+ social media accounts that tracks how policy elites engage with media — broken out by overall affinity, Democratic lean, and Republican lean. That data is not doing enough work today.

Influence Finder turns that dataset into a fast, defensible recommendation that a strategist can hand directly to a client.

2. Who This Is For

Primary user: Early-to-mid-career BPI campaign strategist

Before this tool:

- Pulls a raw data export, filters manually in a spreadsheet
- Cross-references with personal knowledge of the media landscape
- Produces a target list that's hard to defend because the rationale isn't documented
- Spends 2–4 hours on a task that should take 20 minutes

After this tool:

- Enters a campaign brief: client sector, partisan target (D / R / bipartisan), and a media type filter (journalists, outlets, podcasts, etc.)
- Reviews a ranked shortlist of 10–20 accounts with affinity scores, partisan lean, and an elite-vs-general-population index
- Exports a client-ready list — or screenshots the ranked table directly into a deck
- Spends under 15 minutes from launch to deliverable

A day in the life

Scenario

A mid-level strategist is building a media target list for a client making the case for federal voting rights legislation — a campaign where the audience that matters most is Democratic-leaning policy elites in Congress and advocacy circles. She opens Influence Finder, sets lean to Democratic, filters to Journalists, and sorts by Affinity Index. In under ten minutes, she has a ranked list of journalists — names like Karoun Demirjian, Adam Nagourney, and Sewell Chan — ordered not by their raw following but by how disproportionately they over-index with Democratic policy elites specifically. She selects twelve, exports the CSV, and drops the table into the client deck. The rationale is implicit in the ranking and she doesn't have to explain why these names are on the list. The affinity index does that for her.

3. What We're Building — V1

Core interaction

A filter-and-rank interface: the strategist sets a partisan lean target and a media type, and gets a ranked shortlist of targetable accounts, ordered by their affinity index among policy elites.

The primary ranking signal: Affinity Index

Raw reach numbers tell you who's big. The affinity index — a given account's reach among policy elites divided by their reach in the general population — tells you who punches above their weight with the audience that actually matters. An account at 3× affinity is more valuable to a BPI campaign than a bigger account at 1.1×, even if the raw numbers don't show it.

This is the signal that differentiates BPI's recommendation from a list a client could pull from Google.

What V1 does

- Filter by partisan lean: All policy elites / Democratic / Republican / Bipartisan (high on both)
- Filter by media type: journalist, news outlet/website, TV news, newspaper, political podcast — using the `media_tags` field
- Rank results by affinity index (policy elite reach ÷ general population reach), with raw reach and partisan split visible
- Display a clean table: account name, category, affinity index, elite reach %, D/R split, media type
- Allow a strategist to select accounts and export a shortlist

What V1 deliberately does not do

- Issue/beat filtering — the dataset has no journalist topic tags; this belongs in V2 via LLM enrichment or coverage scraping
- Multi-audience support — this version is scoped to the policy elite audience; extension to custom audiences is a V2 build
- Narrative generation — no AI-written rationale copy in V1; the ranked table is the output; auto-generated rationale risks quality issues before we've validated the core workflow
- User accounts or saved searches — V1 is stateless; persistence belongs in V2 once we know what strategists actually want to save

4. Data Notes (as Product Decisions)

A few observations about the dataset that shaped the V1 scope above:

- ~3,300 of 11,500 accounts have a `media_tag` — this is the actionable subset for V1; the remaining rows (entertainment, sports, fashion, etc.) are out of scope for this use case, and we should filter them from the default view
- Some accounts appear to be duplicates (e.g., NY Times, Ny Times, NYT Wirecutter all appear separately) — deduplication logic belongs in V1 as a data prep step, not surfaced to the user as a problem to solve
- No beat/topic data exists at the journalist level — this is the single most significant gap; a strategist building a tech policy campaign can filter to 'journalist' but cannot narrow to 'tech journalists' without this; I'd scope an enrichment sprint for V2
- The affinity index has extreme outliers (accounts with near-zero general population reach score 40× or higher) — we should cap display at a reasonable index ceiling or use a log scale to prevent these from dominating the ranking

5. What Engineering Owns

This brief defines what to build and why. The following are left intentionally open for the engineering lead to decide:

- Tech stack and hosting — no preference expressed; use whatever fits BPI's existing infrastructure
- Deduplication logic — the brief flags that duplicates exist; how to resolve them (fuzzy match, manual curation, data provider fix) is an engineering call
- Affinity index normalization/outlier treatment — the brief flags the problem; the specific approach is an engineering decision
- Export format — CSV is assumed; if the team has a standard export format, use it
- Performance optimization — 3,300 rows is small enough for client-side filtering; no back-end query optimization needed in V1

6. How We'll Know This Is Working

After V1 ships to a pilot group of 3–5 strategists:

Metric	What We're Looking For
Time-to-shortlist	Strategists produce a target list in < 20 minutes (vs. current 2–4 hrs)
List adoption rate	% of shortlisted accounts that make it into actual campaign plans
Client approval rate	Does the client approve the target list without major revisions?
Repeat usage	Are strategists coming back, or opening the tool once and returning to spreadsheets?
Qualitative signal	Do strategists feel the ranking makes sense? Do they trust the affinity index?

The most important signal in the first four weeks is qualitative: are strategists using the output in actual client deliverables? If not, we learn why before building V2.

7. What Comes After V1

The sequence below is intentional — each investment unlocks the next. Beat tagging must precede multi-audience support because issue-level filtering is what makes cross-audience comparison meaningful. Saved searches are only worth building once we know what strategists want to retain, which V1 usage data will tell us. AI rationale is last because it should be trusted, and trust in the ranking comes first.

Phase	Investment · Why this order
V2 · Sprint 1–2	Beat/issue tagging — LLM classification of journalists by coverage area. Unlocks the single most-requested missing filter ('tech journalists who lean Republican'). This is the gap between useful and indispensable.
V2 · Sprint 3–4	Multi-audience support — custom audience segments beyond policy elites, with side-by-side affinity comparison. Requires beat tagging to be meaningful at the segment level.
V2 · Sprint 5	Saved searches + export history — versioned campaign target lists, comparable across clients and time. Only worth building once V1 usage tells us what strategists actually want to save.
V3	AI-generated rationale — one-paragraph plain-English explanation per shortlist, ready to drop into a client memo. Built last, after the ranking is trusted by the team and validated by client outcomes.

8. Open Questions for the Team

- How often does the underlying dataset refresh? Frequency determines whether we need a live data pipeline or a periodic ingest.
- Is there a preference for the export format (CSV, formatted PDF, direct deck integration)?
- Are there existing BPI internal tools this should integrate with, or is a standalone prototype the right V1 form factor?
- Who owns data enrichment (beat/topic tagging)? Is this a BPI Labs initiative or something to request from the data provider?

Appendix A: Data Field Reference

For engineering reference — field names as they appear in the source dataset, their meaning, and how each is used in the V1 product:

Field	Description · Product Use
name	Social account display name. Shown as the primary identifier in the results table. Note: some apparent duplicates exist (e.g., multiple NY Times variants) — requires deduplication in the data prep layer before ingestion.
category	Top-level content category (e.g., 'News Media', 'Politics / Government'). Displayed as a secondary label under the account name. Not used for filtering in V1.
categories_all	Pipe-delimited multi-category string. Not used in V1; retained for potential V2 category expansion.
media_tags	Pipe-delimited type classification (e.g., 'journalist', 'tv_news news_website'). The primary filter field in V1. The first tag in a pipe-delimited string is used as the primary type. Accounts with no media_tag are excluded from the V1 interface entirely.
gen_pop_reach	Decimal fraction (0–1) representing the share of the general population that follows this account. Used as the denominator in the affinity index calculation. Multiply by 100 for a percentage display.
all_reach	Decimal fraction representing the share of all policy elites that follow this account. Primary 'Elite Reach' display column. Denominator for partisan lean classification.
dem_reach	A decimal fraction representing the share of Democratic policy elites that follow this account. Used in partisan lean classification and displayed as the 'Dem' column.
rep_reach	Decimal fraction representing the share of Republican policy elites that follow this account. Used in partisan lean classification and displayed as the 'Rep' column.
affinity_index [derived]	Computed field: $\text{all_reach} \div \text{gen_pop_reach}$. Primary ranking signal. Raw values range from 0.5× to 44.6× in the media-tagged subset; display is capped at 20× in V1 to prevent outliers from distorting the visual bar. The cap is a display decision only — the underlying value is retained for export.
lean [derived]	Computed field: 'democratic' if $\text{dem_reach} \div \text{rep_reach} \geq 2.0\times$; 'republican' if $\leq 0.5\times$; 'bipartisan' otherwise. Used for the partisan lean filter and badge display. Thresholds are a V1 default — recommend reviewing with strategists after pilot.

Appendix B: Dataset Snapshot

Key statistics from the policy elite social account dataset used to inform this brief:

Dimension	Detail
Total accounts	11,560
Accounts with media tags (actionable V1 subset)	~3,300
Media tag breakdown	Journalist (2,172) · Politician (1,172) · News website (404) · TV news (308) · Newspaper (246) · Activist (109) · Political podcast (73) · Other (80+)
Reach signals available	General population reach · All policy elite reach · Democratic policy elite reach · Republican policy elite reach
Affinity index range (media accounts)	1.0× – 44.6× (median ~3–5×; outliers flagged above)
Data gaps	No journalist beat/topic tags · Some apparent duplicate accounts · No temporal dimension (point-in-time snapshot)

— End of Brief —